

Machine Learning

Linear Regression with one variable

By :
Saja Malek Hassan
Maryam Mohamed Ragab
Shahd Ashraf Ramadan

Agenda

Introduction	
Types of Regression	
Problem Statement	
Hypothesis Function	
Cost Function	

Gradient Descent	
Evaluation	
Applications	
Code & Dataset	
Conclusion	

Introduction

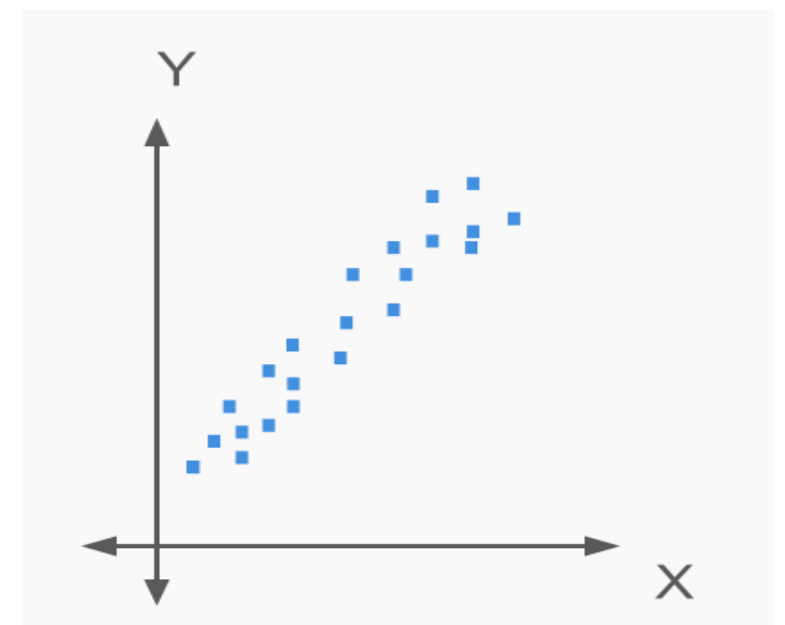
What is regression?

- Supervised learning algorithm that works with **labeled data** .
- Simple model for predicting continuous values .
- It finds the relationship between input features and output.

Linear Regression is a fundamental statistical method used to define **linear relationship** between a dependent and one or more independent variables. It helps predict **outcomes** by fitting a **straight line** to observed data points, making it easy to interpret and apply.

Examples :

- House price from its size.
- salary from years of experience.
- Price of car from Car's attribute.



Types of Regression

Linear Regression

- **Simple Linear Regression** → One independent variable.
- **Multiple Linear Regression** → More than one independent variable.

Logistic Regression

Uses the independent variables to estimate the probability of an event occurring (either 0 or 1)

Polynomial Regression

The regression line is curved rather than straight.

Problem Statement

Scenario: Predicting house price based on its size.

Goal: Find a line that best represents the relationship between size and price, so we can predict :
“If a new house is 100 m², what would its estimated price be?”

House Size (m ²)	Price (\$)
50	150,000
70	200,000
90	250,000
110	300,000

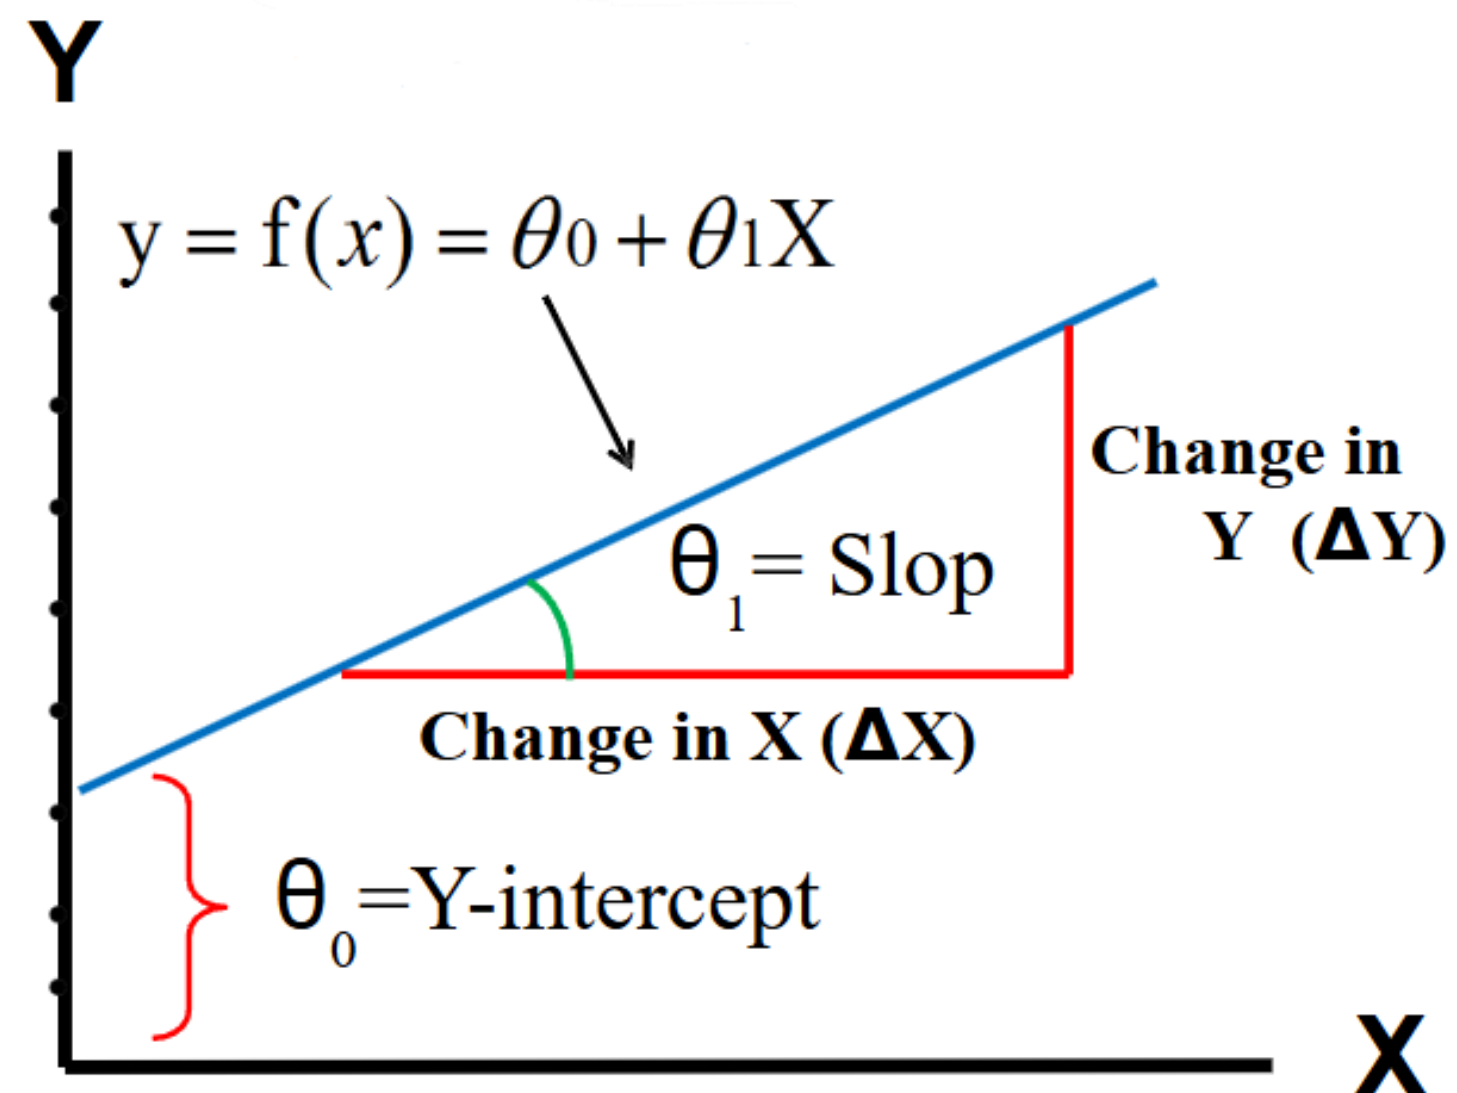
Mathematical Form

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Where:

- $h_{\theta}(x)$: the predicted value (hypothesis)
- θ_0 : intercept (the value of y when $x = 0$)
- θ_1 : slope (how much y changes for each unit change in x)
- x : input variable

Hypothesis Function



Cost Function

The Cost Function measures how well our hypothesis Function fits the data .

It calculate the difference between the predicted value \hat{Y} and the true value Y , also called **Loss Function**.

In Linear Regression, **the Mean Squared Error (MSE)** cost function is employed, which calculates the average of the squared errors between the predicted values \hat{y}_i and the actual values y_i .The purpose is to determine the optimal values for the intercept θ_1 and the coefficient of the input feature θ_2 providing the best-fit line for the given data points. The linear equation expressing this relationship is $\hat{y}_i = \theta_1 + \theta_2 x_i$.

MSE function can be calculated as:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2$$

Where:

- m : number of training examples
- $h_{\theta}(x^{(i)})$:predicted value for the i-th example
- $y^{(i)}$:actual value for the i-th example
- The goal is to **minimize** the cost function $J(\theta_0, \theta_1)$ to find the **best-fitting line**

Cost function visualization

Consider a simple case of hypothesis by setting $\theta_0=0$, then h becomes $h_{\theta}(x)=\theta_1 x$

Each value of θ_1 corresponds to a different hypothesis as it is the **slope** of the line

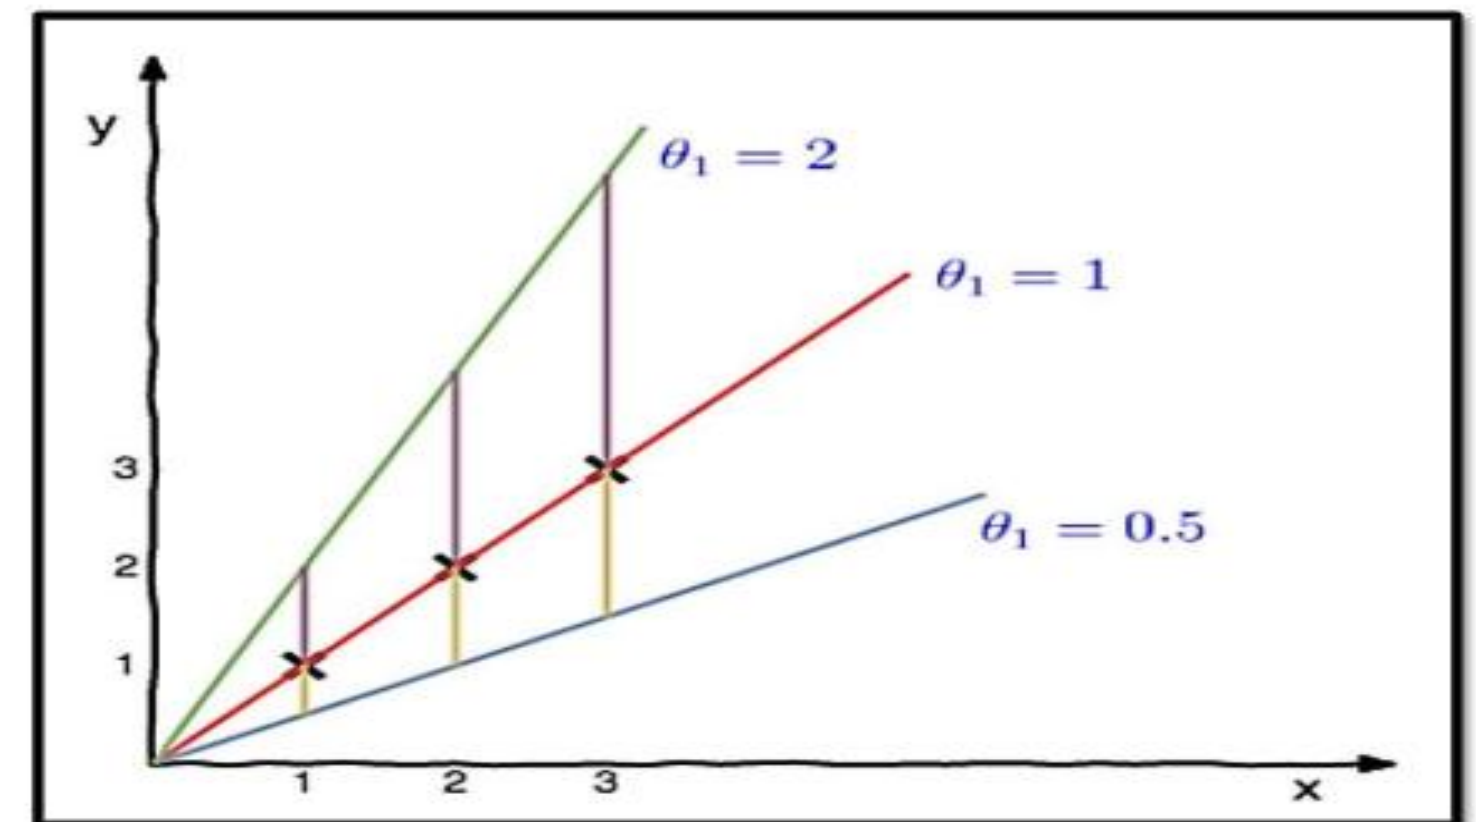
which corresponds to different lines passing through the **origin** as shown in plots below as **y-intercept** i.e. θ_0 is nulled out.

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m \left(\theta_1 x^{(i)} - y^{(i)} \right)^2$$

$$\text{At } \theta_1=2, \quad J(2) = \frac{1}{2 * 3} (1^2 + 2^2 + 3^2) = \frac{14}{6} = 2.33$$

$$\text{At } \theta_1=1, \quad J(1) = \frac{1}{2 * 3} (0^2 + 0^2 + 0^2) = 0$$

$$\text{At } \theta_1=0.5, \quad J(0.5) = \frac{1}{2 * 3} (0.5^2 + 1^2 + 1.5^2) = 0.58$$



Cost function visualization

What is the optimal value of θ_1 that minimizes $J(\theta_1)$?

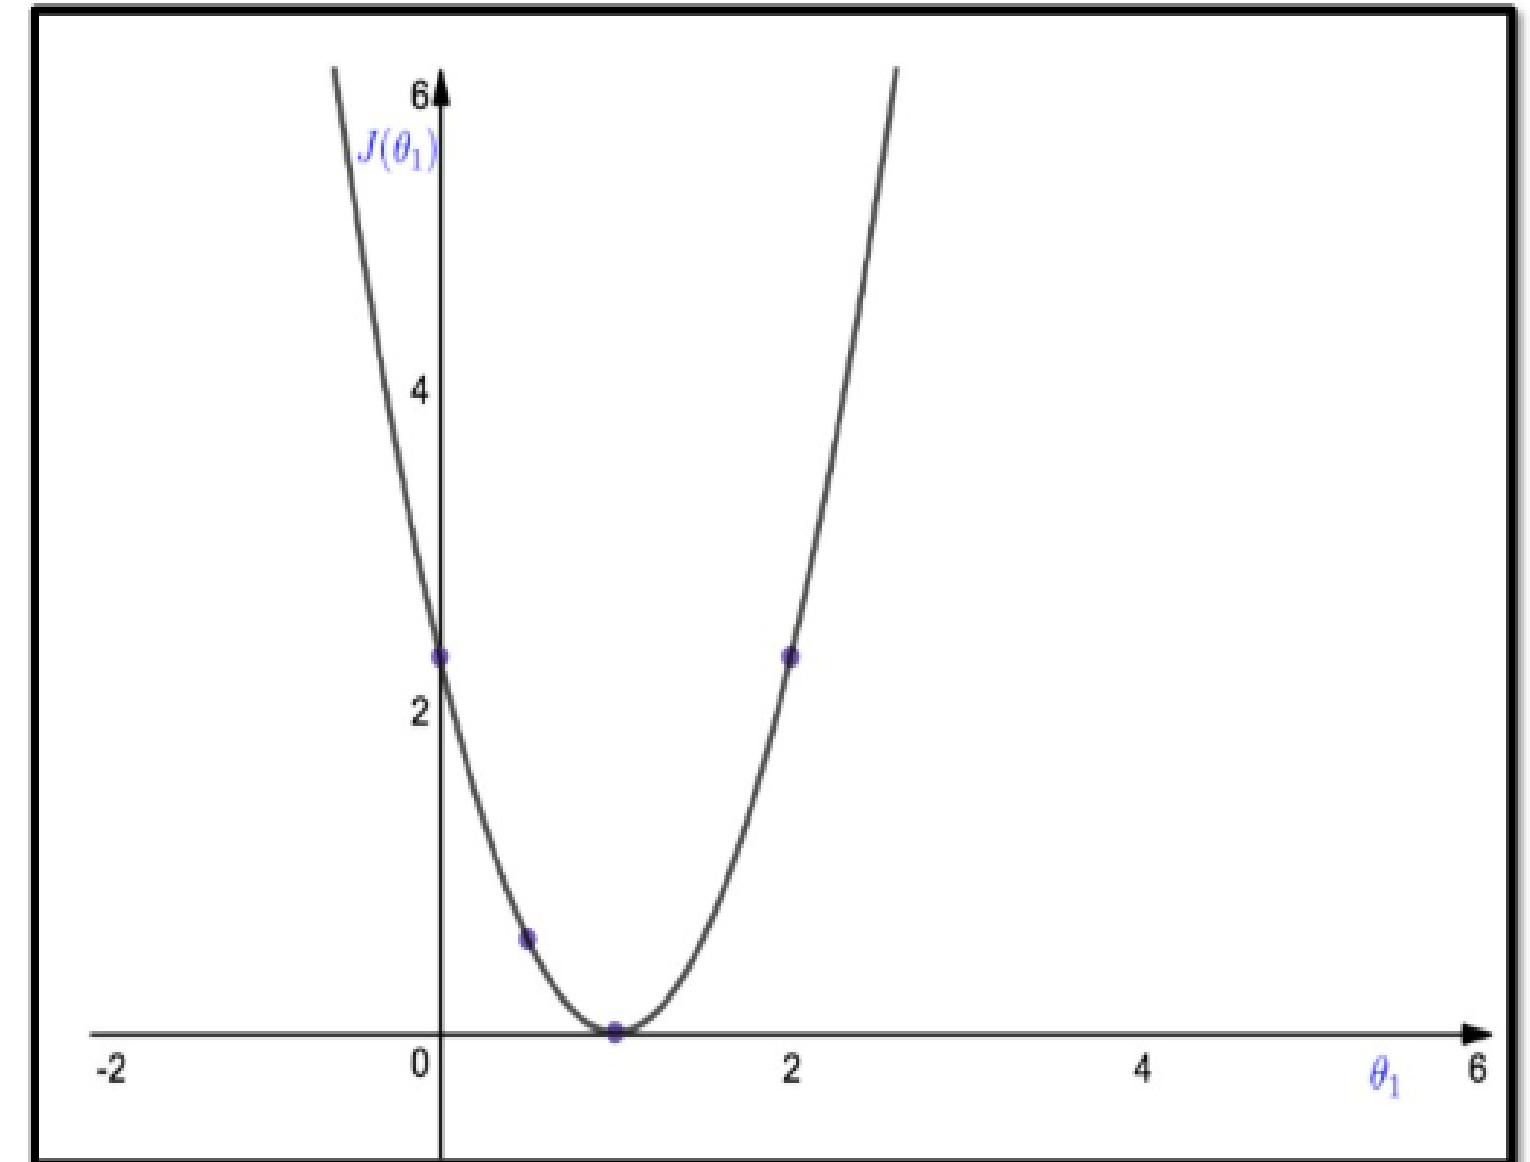
It is clear that best value for $\theta_1 = 1$ as $J(\theta_1) = 0$, which is the minimum.

How to find the best value for θ_1 ?

Plotting ?? Not practical specially in high dimensions?

The solution :

1. Analytical solution: not applicable for large datasets
2. Numerical solution: ex: Gradient descent .



Cost Function

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Parameters: θ_0, θ_1

Cost Function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Goal: $\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$

Gradient Descent

Gradient Descent

- ◆ Is an Optimization algorithm used to minimize the cost function.
- ◆ It updates the parameters θ_0 , θ_1 step by step in the direction that reduce the errors.
- ◆ The process continues until the cost function reaches its minimum value.
- ◆ Iterative Solution not only in Linear Regression , it's actually used all over the place in machine learning .

Gradient Descent

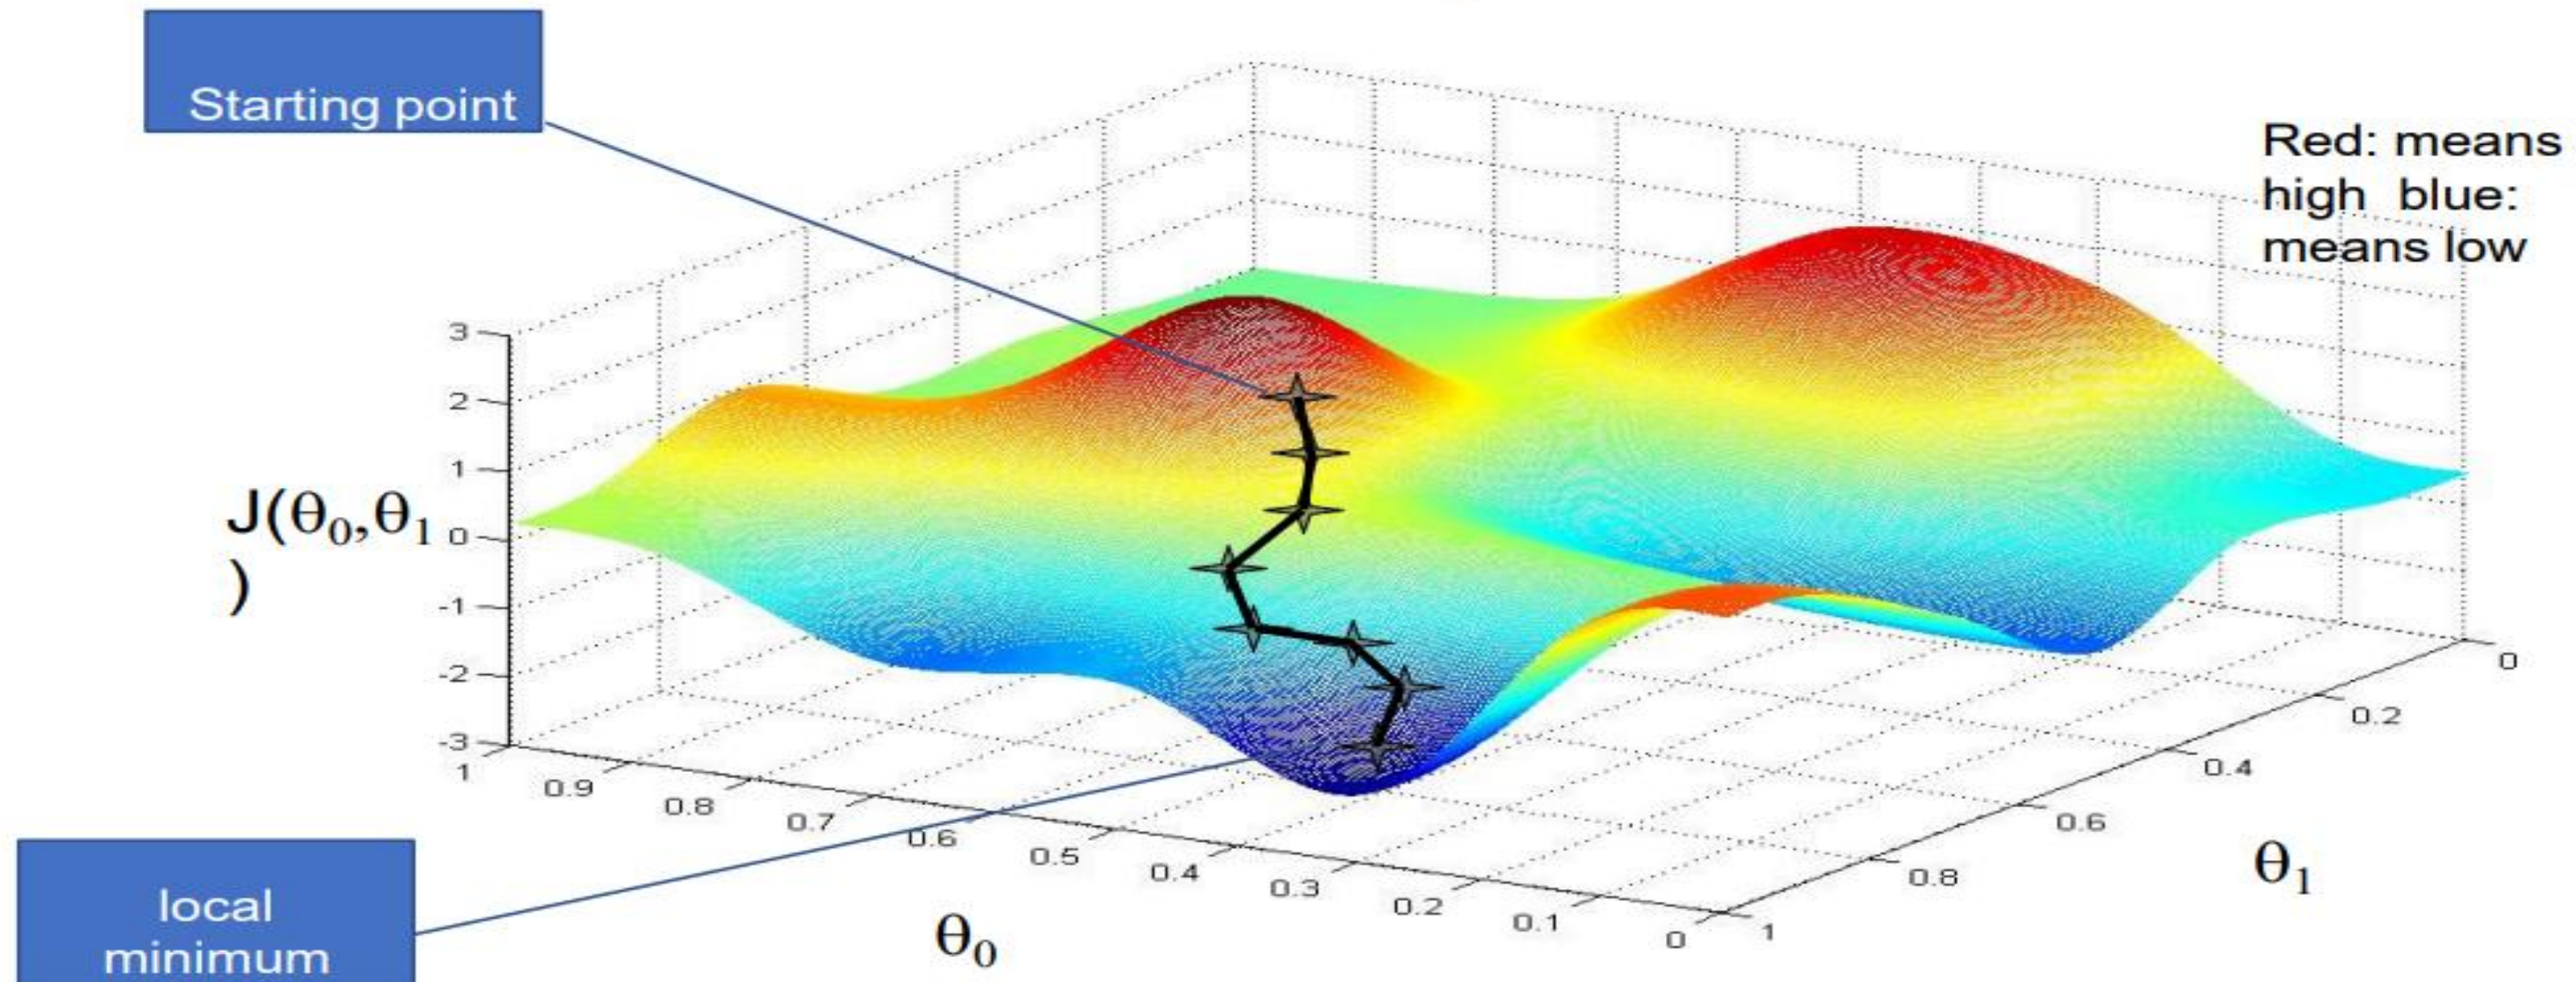
Have some function $J(\theta_0, \theta_1)$

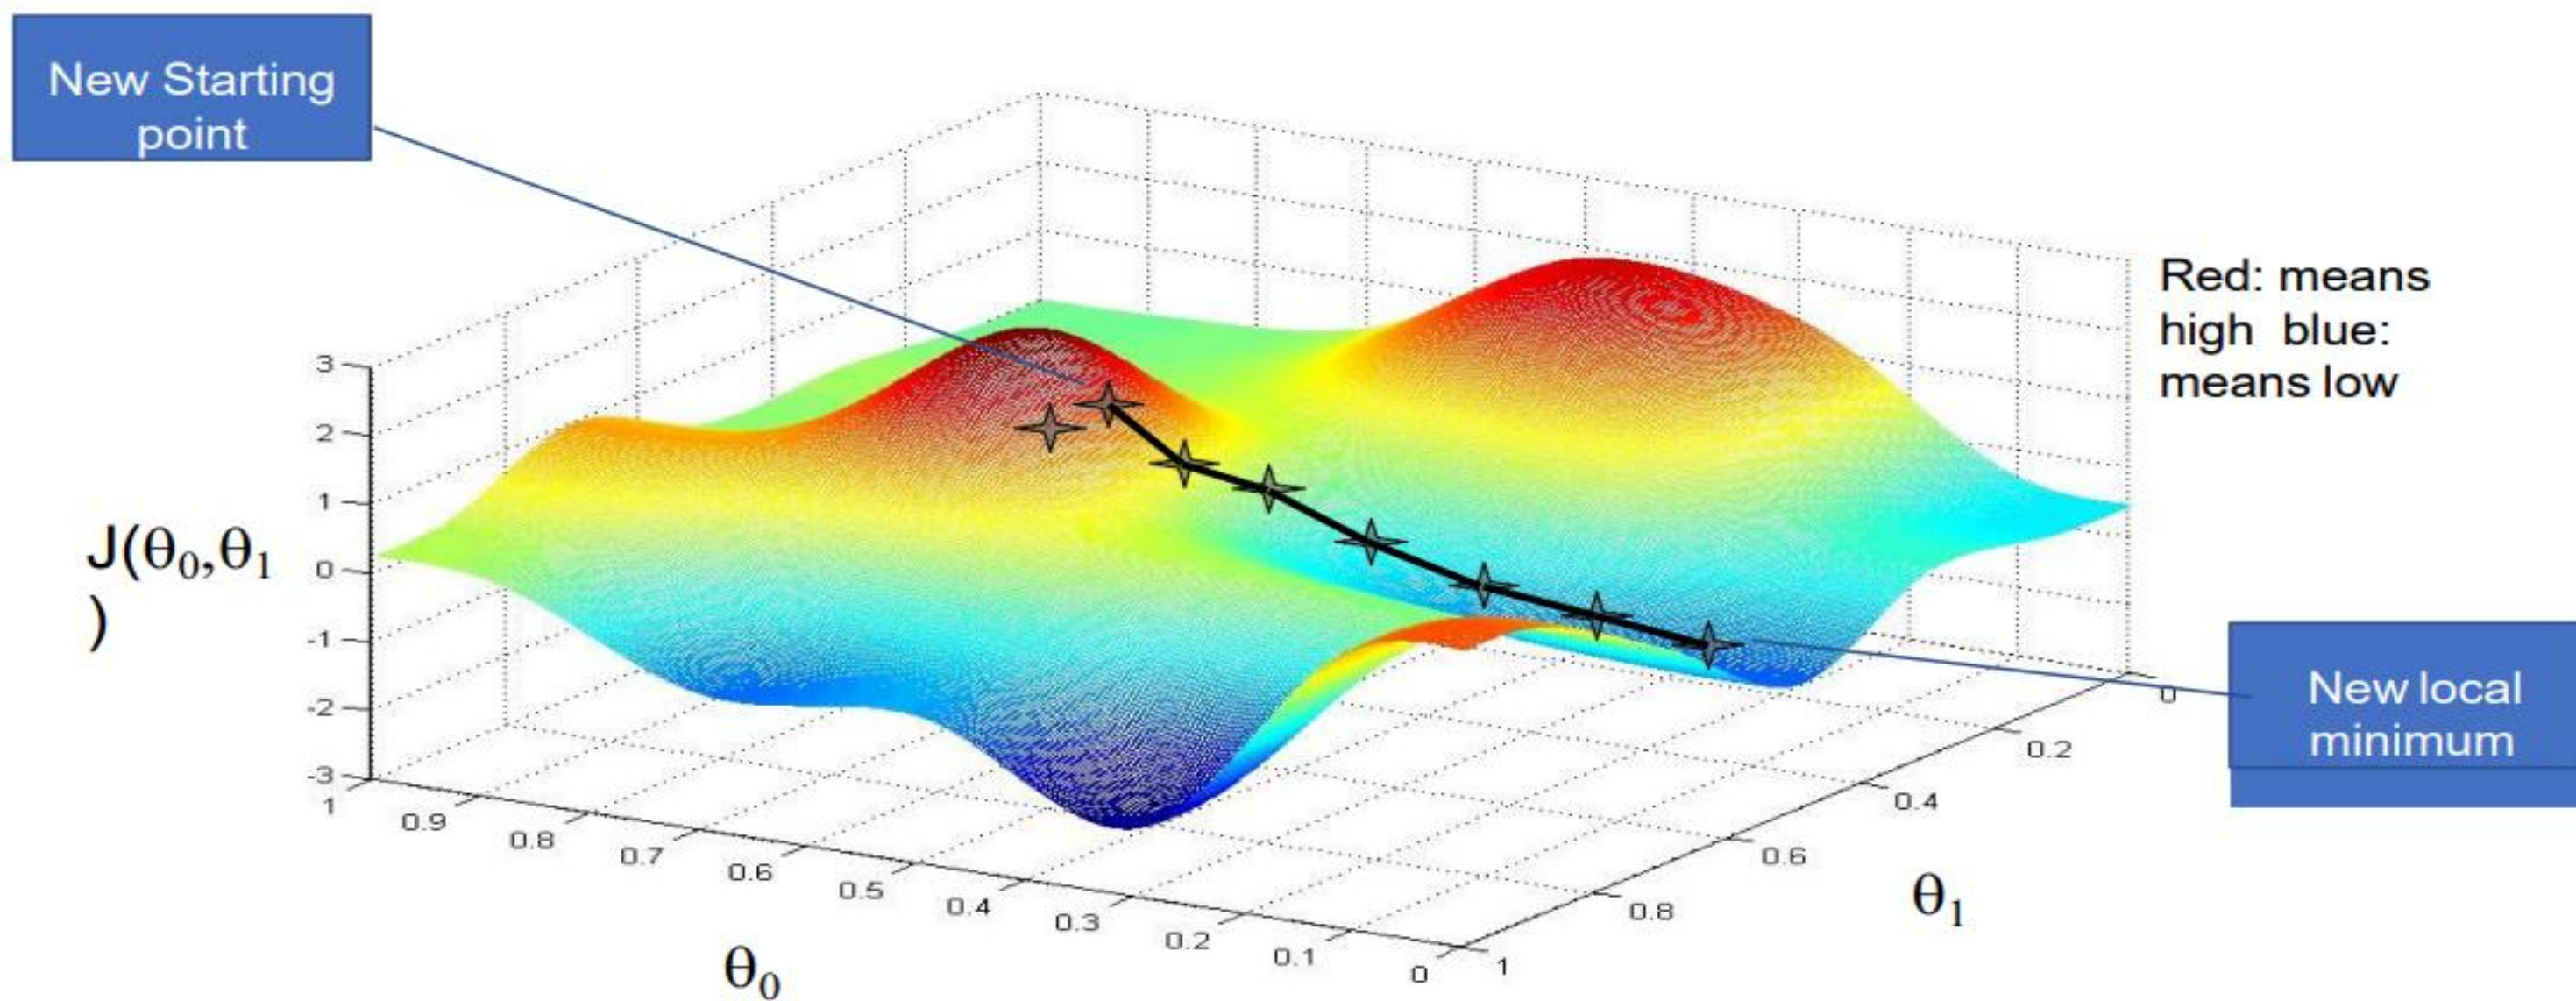
Want $\min J(\theta_0, \theta_1)$

Outline :

- start with some θ_0, θ_1
- Keep changing θ_0, θ_1 *to reduce* $J(\theta_0, \theta_1)$
until we hopefully end up at a minimum .

Imagine that this is a landscape of grassy park, and you want to go to the lowest point in the park as rapidly as possible





With different starting point

Gradient Descent Algorithm

repeat until convergence $\{\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \forall j \in \{0, 1\}\}$

- Where

- $:=$ is the assignment operator
- α is the **learning rate** which basically defines how big the steps are during the descent
- $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ is the **partial derivative** term
- $j = 0, 1$ represents the **feature index number**

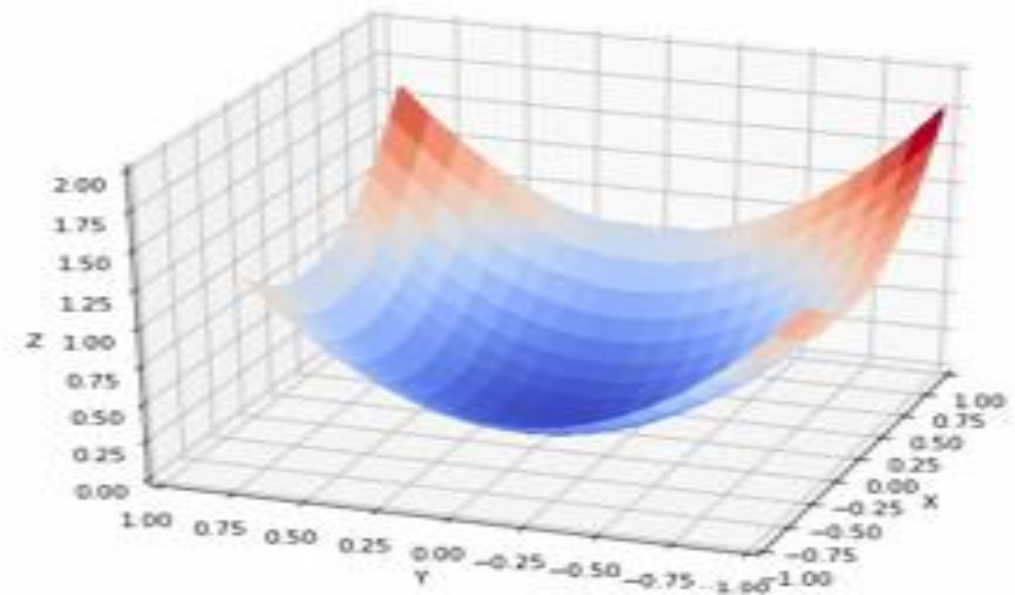
Also the parameters should be **updated simulatenously**, i.e. ,

$$temp_0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$temp_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

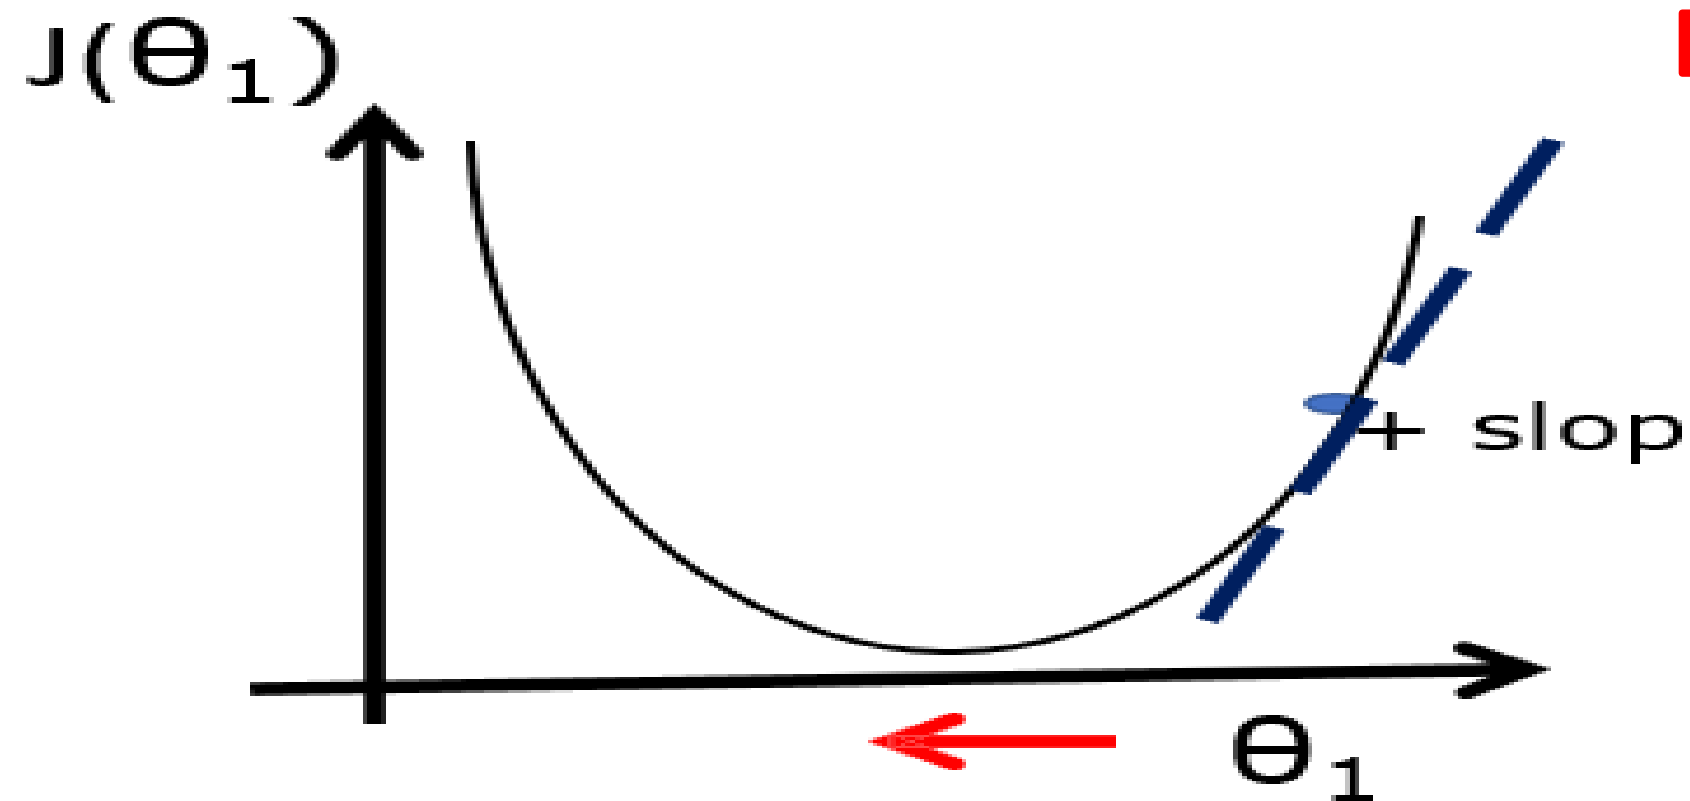
$$\theta_0 := temp_0$$

$$\theta_1 := temp_1$$



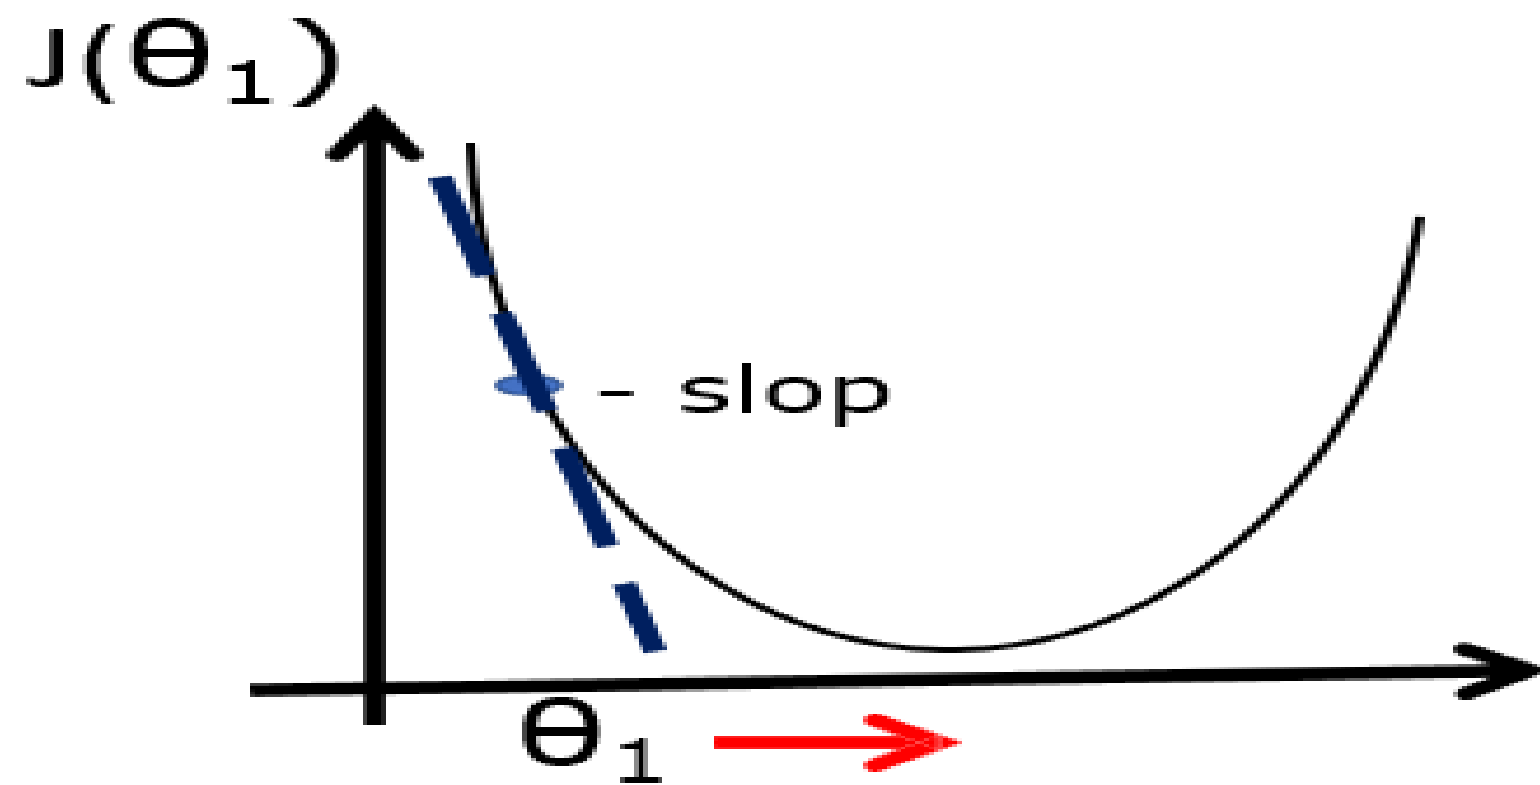
Gradient Descent Algorithm

Example



$$\theta_1 = \theta_1 - \alpha \frac{d}{d\theta_1} j(\theta_1)$$

$$\theta_1 = \theta_1 - \alpha (+ve) \quad \downarrow$$



$$\theta_1 = \theta_1 - \alpha (-ve) \quad \uparrow$$

Gradient Descent For a linear Regression

Gradient Descent Algorithm

Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
 (for $j = 1$ and $j = 0$)
}

Linear Regression Model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient Descent Algorithm

$$\frac{d}{d\theta_j} j(\theta_0, \theta_1) = \frac{d}{d\theta_j} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - Y_i)^2$$

$$\frac{d}{d\theta_j} j(\theta_0, \theta_1) = \frac{d}{d\theta_j} \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 (x_i) - Y_i)^2$$

$$j=0: \frac{d}{d\theta_0} j(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x_i) - Y_i)$$

$$j=1: \frac{d}{d\theta_1} j(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x_i) - Y_i) \bullet x_i$$

Gradient Descent Algorithm

repeat until convergence {

$$\left. \begin{aligned} \theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \\ \theta_1 &:= \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)} \end{aligned} \right\} \begin{array}{l} \text{update} \\ \theta_0 \text{ and } \theta_1 \\ \text{simultaneously} \end{array}$$

}

Gradient Descent Algorithm

“Batch” Gradient descent

Batch : Each step of gradient descent uses all the training examples .

repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

update
 θ_0 and θ_1
simultaneously

}

Gradient Descent Algorithm

◆ Batch Learning:

Use all training examples to update the parameters in each gradient descent algorithm.

◆ Online Learning:

Use each individual training example to update the parameters in each gradient descent algorithm.

Evaluation

Evaluation

Model Evaluation Metrics

1. Mean Squared Error (MSE)

$$MSE = \frac{1}{m} \sum_{i=1}^m \left(\widehat{y^{(i)}} - y^{(i)} \right)^2$$

- Used to measure the average squared difference between predicted and actual values. Best when you want to penalize larger errors more heavily.
- A smaller MSE value means better model performance.

2. Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{\frac{1}{m} \sum_{i=1}^m \left(\widehat{y^{(i)}} - y^{(i)} \right)^2}$$

- *Provides the error in the same unit as the target variable (Y).*
- *Easier to interpret compared to MSE.*
- *Commonly used when the scale of the prediction error is important.*

Evaluation

Model Evaluation Metrics

3. R^2 Score (Coefficient of Determination)

$$R^2 = 1 - \frac{\sum_{i=1}^m \left(y^{(i)} - \widehat{y}^{(i)} \right)^2}{\sum_{i=1}^m (y^{(i)} - \bar{y})^2}$$

Measures how much of the variance in the target variable is explained by the model.

Perfect fit $\rightarrow R^2 = 1$

Model explains none of the variance $\rightarrow R^2 = 0$

Higher R^2 indicates better model performance.

Applications



1. Business & Economics

Predicting sales , profits ,or market tends based on advertising spend,pricing , demand .



2. Helthcare

Predicting disease progression or patient recovery time from medical .



3. Education

Predicting student performance based on study hours , attendance , and previous scores .

Applications



4. Engineering

Modeling relationships between Variables , such as temperature vs.pressure , or stress vs. strain .



5. Real Estate

Estimating house prices from features such as , size , number of room , and location .



6. Finance

Forecasting stock prices , interest rates or credit risk using historical data .

Code & Dataset

Access the implementation and data here:

https://drive.google.com/drive/folders/1b1BaCqWzDtVxlc-iBU0BoP3nZBIS_NJQ?usp=sharing

Conclusion

- ◆ Linear Regression is one of the simplest and most important algorithm in machine learning .
- ◆ It helps to find the relationship between input and output variables and predict continuous values .
- ◆ The model works by fitting a straight line that minimizes the cost function (error) using Gradient descent .

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Cost Function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Thank You